

Combinatory Categorical Grammar and Link Grammar are Equivalent

Linas Vepstas

13 July 2022

Abstract

This is a short, semi-formal note explaining how **Combinatory Categorical Grammar** (CCG) and **Link Grammar** (LG) are equivalent. It covers some basic ideas from proof theory, type theory, and the “sexuality” of type combinators. The key idea that exposed is that type theory must be combined with connector sexuality in order to get a fully general framework encompassing proof theory, inference in logic and linguistics.

A Question posed on a Discord chat channel

@Adam Vandervorst asks: *Does anyone here know about https://en.wikipedia.org/wiki/Combinatory_categorical_grammar?*

(from Wikipedia) Combinatory categorical grammar (CCG) is an efficiently parsable, yet linguistically expressive grammar formalism. It has a transparent interface between surface syntax and underlying semantic representation, including predicate–argument structure, quantification and information structure. The formalism generates constituency-based structures (as...

I talked to its inventor last week (who has since moved on to do learning in linguistics) and it was really interesting — Today at 9:31 AM

The Nature of Grammar

To open, there’s this thing about grammars that you should know.

As far as I can tell, all of the different (formal) grammar formalisms are inter-convertible into one-another, by purely algorithmic means. That is, given the collection of symbols and rules that are used to define one formalism (*e.g.* constituency grammars, CG) one can convert that into a different formalism (*e.g.* dependency grammars, DG) by applying a purely automatic transformation on the grammar specifications. No hand-waving is required, nor any metaphysics: a machine can convert DG into CG and *vice versa*, and that machine is rather simple.

Somewhere out there is a nice paper that explains how to convert between DG and CG and back. It provides a simple algo to do this. Sadly, I have misplaced the reference.


CCG is Equivalent to LG

I looked at CCG many years ago, and from what I could tell, for each and every CCG compound type, one has an equivalent LG link type, and *vice versa*. For example, the compound type NP/N is the same thing as the LG D^+ link (determiner) type and $(S/NP)/NP$ is the just the LG $S- \& O^+$ (verb taking subject and object) and one can march down the list this way. The goal of this PDF is to make the above statement precise.

At first, it's mildly confusing, because it seems like the compound type NP/N might be encoding some kind of structure that the single-letter, single-type D^+ is not ... but, actually, no, that is incorrect. The CCG notation is not actually "more atomic" or "more compositional" than the LG notation. To understand this, one must slightly shift one's point-of-view.

Jigsaw Pieces

Recall how I talk about "jigsaw pieces" all the time? Some example LG jigsaw pieces:



The above diagram is taken from the original 1991 paper presenting Link Grammar.^[1] Now, let's look at that CCG Wikipedia article. You can find this inference rule:

$$\frac{\alpha : X/Y \quad \beta : Y}{\alpha\beta : X} >$$

This says that (roughly speaking) "if you have a jigsaw called alpha and it has connector of type X on left and type Y on right, and if you have jigsaw beta with a connector Y , you can connect the two Y 's together, to yield a combined jigsaw alpha beta having only one unconnected connector X ."

Let's now try to be more precise. This inference rule uses conventional proof-theory style notation. The horizontal line defines the inference to be done. Above the line are the inputs, below the line are the outputs. The Greek letters α, β are terms, and the Latin letters X, Y are types. The colon indicates a term-type pairing, so that $\beta : Y$ is a term of type Y . The slash $/$ and the backslash \backslash are **type constructors**, so that X, Y are simple types, and X/Y is a compound type, constructed from the two simpler types. The $>$ is just a label for the rule; it has no syntactic role.

The CCG Wikipedia article calls these inference rules “combinators”. Above is one “application combinator”; there is also a second rule:

$$\frac{\beta : Y \quad \alpha : X \backslash Y}{\beta \alpha : X} <$$

Lets rewrite these two rewrite rules in LG notation. They would be

$$\frac{\alpha : (X- \ \& \ Y+) \quad \beta : Y-}{\alpha \beta : X-} >$$


and

$$\frac{\beta : Y+ \quad \alpha : (X+ \ \& \ Y-)}{\beta \alpha : X+} <$$

Here, the X, Y are the LG types, called “link types” in the literature. The $Y+$ and $Y-$ are called “connectors”: they are jigsaw-puzzle-piece tabs, as-yet unconnected. When they do connect, they are called a “link”, and thus the name “Link Grammar”. The $+$ and $-$ are the “connector directions”: they specify in which direction a connector can connect: to the right or to the left.

The ampersand $\&$ is a “kind of” type constructor. Given two connectors, say, $X-$ and $Y+$ it creates a new type (more precisely, a “jigsaw”) $X- \ \& \ Y+$. This can be made more precise, in an upcoming section.

So what are these two inference rules really saying? Well, its almost trivial: they’re just saying “connectable connectors can connect, if the connector types are identical, and the sexuality of the connectors is opposite.” Let’s cement the obvious. Here’s the first combinator, using the same diagrammatic representation as in the original 1991 Link Grammar paper:



What is this picture saying? The obvious: when you combine terms α and β the result is a single term $\alpha\beta$ and it is convenient to not draw, to ignore, to pretend that the link Y joining these two pieces as disappeared. In other words, a partially-assembled jigsaw puzzle behaves exactly like a single jigsaw piece.

The CCG Composition Combinators

For completeness, the remaining CCG combinators should be treated as well. Here’s a side-by-side Rosetta Stone of the two composition combinators.

CCG	LG
$\frac{\alpha:X/Y \quad \beta:Y/Z}{\alpha\beta:X/Z} B_{>}$	$\frac{\alpha:X- \& Y+ \quad \beta:Y- \& Z+}{\alpha\beta:X- \& Z+} B_{>}$
$\frac{\beta:Y\backslash Z \quad \alpha:X\backslash Y}{\beta\alpha:X\backslash Z} B_{<}$	$\frac{\beta:Y+ \& Z- \quad \alpha:X+ \& Y-}{\beta\alpha:X+ \& Z-} B_{<}$

Clearly, they just specify how to connect compound connectors.

The CCG Type-raising Combinators

The last pair of combinators are the type-raising combinators. These are

CCG	LG
$\frac{\alpha:X}{\alpha:T/(T\backslash X)} T_{>}$	$\frac{\alpha:X-}{\alpha:T- \& T+ \& X-} T_{>}$
$\frac{\alpha:X}{\alpha:T\backslash(T/X)} T_{<}$	$\frac{\alpha:X+}{\alpha:T+ \& T- \& X+} T_{<}$

The interpretation of these two rules is that, given a single (assembled) jigsaw X , cut it into two (disassemble it) such that the new connectors are of type T (and they are no longer connected).

There seems to be a slight awkwardness, as the earlier combinators could be easily understood by thinking only about simple types. By contrast, the CCG requires a more complex explanation:

“The type-raising combinators, often denoted as $T_{>}$ for forward type-raising and $T_{<}$ for backward type-raising, take argument types (usually primitive types) to functor types, which take as their argument the functors that, before type-raising, would have taken them as arguments.”

Phew. That’s a mouthful, when all that is really being said is “disconnect” or “cut into pieces”. A bit more on why and how this extra complexity arises, shortly below.

An Sloppy Example


The Wikipedia article includes an example of two different proofs (two different derivation trees) of the same sentence. Here's one:

$$\frac{\frac{\frac{\text{the}}{NP/N}}{NP} > \frac{\frac{\text{dog}}{N}}{N} > \frac{\frac{\frac{\text{bit}}{(S\backslash NP)/NP}}{S\backslash NP} > \frac{\frac{\text{John}}{NP}}{NP} >}{S} <$$


If we are sloppy and uncaredful ***, we find that the translated LG derivation rules:

$$\frac{\frac{\frac{\text{the}}{NP+ \& N+}}{NP+} > \frac{\frac{\frac{\text{dog}}{N-}}{N-} > \frac{\frac{\frac{\text{bit}}{S+ \& NP- \& NP+}}{S+ \& NP-} > \frac{\frac{\text{John}}{NP-}}{NP-} >}{S+} <$$

This is perhaps hard to read? The conventional LG notation for this derivation would be:



where an extra jigsaw piece RIGHT-WALL: S- was introduced, so as to keep all connectors fully connected. The above works. It is not the preferred LG parse for the current English language dictionary. That would be:



The link types are obviously more complex. Note also the present of a cycle (the triangle, whose edges are WV, Wd and Ss*s.) Note the presence of several directed connectors. The complex upper-case/lower-case link types are an example of “sexuality”; see next section.

A Less Sloppy Example

*** Wait, what? Sloppy and uncaredful? If we are careful, and don't gloss any plus and

minus signs, then the following derivation results:

$$\frac{\frac{\text{the}}{NP- \& N+} \quad \frac{\text{dog}}{N-}}{NP-} > \frac{\frac{\text{bit}}{S+ \& NP- \& NP+} \quad \frac{\text{John}}{NP-}}{S+ \& NP-} > \frac{}{NP- \& S+ \& NP-} \text{ fail !!}$$

This reveals a bug in the Wikipedia article derivation. It should have been:

$$\frac{\frac{\text{the}}{N} \quad \frac{\text{dog}}{NP \setminus N}}{NP} < \frac{\frac{\text{bit}}{(S \setminus NP) / NP} \quad \frac{\text{John}}{NP}}{S \setminus NP} < \frac{}{S} <$$

This provides a hint as to why LG might actually be better than CCG: it's easier to spot bugs. We live in an era of compilers and debuggers; yet hand-writing expressions is error prone.

The CCG article also gives an alternative derivation for the sentence, but it is also buggy. The article currently states:

$$\frac{\frac{\frac{\text{the}}{NP/N} \quad \frac{\text{dog}}{N}}{NP} > \frac{}{S/(S \setminus NP)} T > \frac{\frac{\text{bit}}{(S \setminus NP) / NP} \quad \frac{\text{John}}{NP}}{S/NP} B > \frac{}{S} >$$

The naive translation of the above to LG derivations reveals two bugs in the above:

$$\frac{\frac{\frac{\text{the}}{NP- \& N+} \quad \frac{\text{dog}}{N-}}{NP-} > \frac{}{S- \& S+ \& NP-} T > \frac{\frac{\text{bit}}{S+ \& NP- \& NP+} \quad \frac{\text{John}}{NP-}}{S- \& S- \& NP- \& S+ \& NP-} B > \frac{}{XXX} \text{ fail !}$$

It should have been this:

$$\frac{\frac{\frac{\text{the}}{N} \quad \frac{\text{dog}}{NP \setminus N}}{NP} < \frac{}{S \setminus (S/NP)} T < \frac{\frac{\text{bit}}{(S \setminus NP) / NP} \quad \frac{\text{John}}{NP}}{S/NP} B > \frac{}{S} >$$

This corrected form then translates as:

$$\frac{\frac{\frac{\text{the}}{N+} \quad \frac{\text{dog}}{NP+ \& N-}}{NP+} < \frac{}{S+ \& S- \& NP+} T < \frac{\frac{\text{bit}}{S+ \& NP- \& NP+} \quad \frac{\text{John}}{NP-}}{S+ \& NP+} B > \frac{}{S+} >$$

One of the two bugs is the same as the earlier one. The second bug is more subtle: it was a misapplication of the $B_{>}$ rule. The two premises of the $B_{>}$ rule must necessarily have inverted polarities; the Wikipedia article gives them the same polarity, and then cancels them out. To be explicit: the compound type $(S \backslash NP)$ cannot be canceled against another $(S \backslash NP)$ in the $B_{>}$ rule. It must be canceled against (S / NP) .

If this analysis seems incorrect, muddled or confused, or anchored in in misinterpretation, that's OK. Don't give up; it's complicated. A deeper analysis will be presented in the section on connector sexuality. In short, the default presentation of CCG assumes mono-sexual types (types without the $+/-$ directional markup), but this will be seen to not be enough.

The root cause of both of these bugs was a failure to attend the polarity that is implied by the type constructors $/$ and \backslash . These type constructors build compound types with an implicit polarity; the failure to write it down leads to interpretational issues. These bugs can only be resolved by taking care to distinguish between types and sexualities (polarities, here, since the sexualities here are heterosexual.) More on sexuality, shortly. If the above analysis seems incorrect, muddled or confused, or anchored in in misinterpretation, please keep reading.

There is yet another infelicity in the Wikipedia article. It currently states:

The sentence "the dog bit John" has a number of different possible proofs. Below are a few of them. The variety of proofs demonstrates the fact that in CCG, sentences don't have a single structure, as in other models of grammar.

This is misleading. Two different derivation trees are presented. The ultimate parse is identical. This phenomenon is commonly treated in textbooks on proof theory: two different proofs have proof trees that appear to be different, but can be rearranged by homotopic deformations into one-another. That is, there is a Scott-continuous deformation, referring to the Scott topology that conventionally applied to proofs/programs. How can continuous transformations be spotted? Next section!

Conclusion

In conclusion: CCG is equivalent to LG. The inference rules of CCG are merely rules for how to join together connectors. Two rules connect simple types to compound types; two more rules connect compound types, and the final two rules show how to disassemble connections (equivalently, to create unconnected pairs).

It should be clear that CCG uses a far more awkward notation (the proof-theoretical inference-bar notation). Awkwardness matters, because concepts like link-crossing and Dick Hudson's "landmark transitivity" becomes very hard to talk about in CCG.


Proof Theory

Although the presentation above focused on CCG, and LG, the concept of inference rules as being certain peculiar kinds of rewrite rules is not new. Lets take a look at the

“standard form” of an **inference rule**, taken from Wikipedia:

Premise #1
 Premise #2
 ...
 Premise #n
 —————
 Conclusion

This is, oddly enough, just another jigsaw. Let’s be painfully clear, by actually drawing it:



The X_k are the premises, the Y is the conclusion. These are drawn as it they’re typed; the type assignments are just illustrative; what matters in this picture are the directionalities: there are n inputs and one output. Structurally, this has the form of a lambda combinator, having n inputs ... in practical applications, inference rules behave as if they were lambdas. The central point being made here is that the input-to-output connections are heterosexual. Premises cannot be “plugged into” premises; conclusions cannot be “plugged into” conclusions. There only one possible direction: only conclusions can be plugged into premises.

All proof-theoretical inference rules are always jigsaw pieces. All of them, without any exceptions. This holds for any type of logic: classical, predicate, intuitionist, modal, linear logic. This observation is "trivial" because its effectively just a notational thing.

Alternative notations used to write inference rules, however, is interesting. One common form appearing in many computer-science settings is

$$X_1 \wedge X_2 \wedge \cdots \wedge X_n \rightarrow Y$$

The LG notation for this is

$$X_1- \ \& \ X_2- \ \& \ \cdots \ \& \ X_n- \ \& \ Y+$$

In tensor algebras, one would write

$$X_1^* \otimes X_2^* \otimes \cdots \otimes X_n^* \otimes Y$$

where the $*$ denotes the (contravariant) dual. In index notation, this would be written as

$$T_{\alpha\beta\cdots\mu}{}^{\nu}$$

In quantum mechanics, one uses the bra-ket notation:

$$|X_1\rangle \otimes |X_2\rangle \otimes \cdots \otimes |X_n\rangle \langle Y|$$

One must be careful, though; the last three forms can be beguilingly misleading. Tensor algebras are (dagger) symmetric, and thus have only one type constructor. In linguistics, there are two type constructors, which go to the left and the right, because the left-right distinction matters in linguistics. This is a source of confusion that hasn't been (in my mind) fully and clearly resolved. There's a further note on this at the end.

Connector Sexuality

In CCG, there are two "type raising combinators" \backslash and $/$ because in linguistics, word order matters. Nouns appearing to the left of a verb are subjects; nouns appearing on the right are objects. Link Grammar accomplishes the left-right distinction with the $+$ and $-$ connector directions. This is, in general, sufficient for linearly-ordered sequences of words.

The rules for joining together LG connectors state that $+$ can only be attached to $-$. One can never attach $+$ to $+$ or $-$ to $-$. In this sense, the connection rules are heterosexual. This is also the conventional mechanism for lambda calculus, and of function calls: one can plug earlier outputs into new inputs. One can take a number 42 and plug it into $f(x)$ to get $f(42)$ but one cannot plug $f(x)$ into 42. Nor 42 into 42, for that matter. Function calls are also heterosexual (and almost always typed, except for simply-typed lambda calculus).

Mono-sexual connectors are those for which there is only one connector type. It can be denoted simply by $*$, or not at all (by just dropping the concept). In a monotyped system, all relations are necessarily homosexual, as there is only one sex.

Jigsaws can in general have monosexual connectors, or trisexual connectors, or other arbitrarily complex rules.

Trisexuality

An example of trisexual connectors would be the set

$$\{Aa, Ab, Ac, Ba, Bb, Bc\}$$

with the connection rules that upper-case letters must match, while lower-case letters must be different. In this case, there are three sexes a, b, c instead of two $+-$, but the rules still demand heterosexuality between the connector "directions".

One can enliven the situation by introducing $*$ as a direction wild-card. Thus, $*$ can mate with $*$, or with any of the sexes a, b, c . So for example, Ab can attach to $A*$, or any of the other A 's; just not to another Ab . Likewise $B*$ cannot attach to any of the A 's, because the uppercase letters denote type, and you cannot mix these.

The fundamental need for connectors

We now come to perhaps the most subtle point of this. It's subtle because its blaringly, forehead-slappingly obvious. It's so obvious that, in fact, it will shoot right by, if you are not paying attention.

It is this: in almost all conventional, day-to-day usage of types, when someone says “this is a type”, half the time, they really mean “this is a connector”. Connectors are implicitly present almost everywhere; their use is rampant, and the concept of “direction” is never mentioned, because it is almost always obvious from context. One could say that type theory and computer programming suffer from “systemic hetero-sexism” or “normative heterosexuality”.

Consider programming in C, C++ or Java. Three basic types are `int`, `float` and `string`. Duhh. Function calls have “signatures”, *e.g.* `int func(int x)`. What is the number 42? Obviously, it’s an `int`, and obviously you can plug it into `int func(int x)`, so that `func(42)` is syntactically valid (in C, C++, Java) but `42(func)` is obviously syntactic nonsense. No one ever needs to explain this.

It would be strange and bizarre to explain that 42 is actually a “connector”, having type `int` and direction “output”. Likewise, in `func(int x)`, the `x` is actually connector. It obviously has the type `int`, but it also has the direction of “input”. The connector rules that connections can only be heterosexual are obvious: you can connect an output to an input, and that is it. You cannot connect two inputs, you cannot connect two outputs. Duhh. Any dummy knows this.

Now it is time to slap one’s forehead. About half the time when someone says something is of type X what they really mean is that something is a connector, and that connector has a type of X and a direction of either “input” or “output”, which is always obvious from context.

Normative Heterosexuality

The reason for my belaboring this “normative heterosexuality” is that sometimes, it gets you into trouble. When CCG writes the inference rule

$$\frac{\alpha : X/Y \quad \beta : Y}{\alpha\beta : X} >$$

the types X and Y were implicitly mono-sexual. They were taken to have no directional information, and all left-right distinctions in the grammar emerged from the two type constructors $/$ and \backslash . Superficially, this seems all fine and correct, although ambiguous associative situations arise, which can be resolved by using parenthesis. Thus, associative expressions such as $X \backslash (Y/Z)$ and the algebra of CCG types is a non-associative algebra (the locations of the parenthesis matter).

In fact, the mono-sexed types, when used with the two combinators and with the parenthesis, provides a golden path to hell. This doesn’t become apparent until one starts tripping over buggy expressions. The two example sentences contained three bugs, grand total. These bugs were not visible until the placeholders X and Y were reinterpreted as connectors (with types X and Y), and the previously implicit directional attachments were made explicitly visible with $+$ and $-$.

Constituency and Dependency

Wait! It gets even worse!

The reason for the first bug was a slavish adherence to the conventions of old-fashioned constituency grammar. The inherited tradition is that N denotes a noun, and NP is a noun phrase. When one writes “*The dog bit John*”, it is clear that “*the dog*” is an NP, and it is also clear that “*the*” is not N, and that “*dog*” is N. Thus, one is forced into assigning NP/N to the determiner. But this is an error!

The markup NP/N is saying that the word “*the*” is a noun-phrase, and it’s just missing a noun before it becomes a complete NP. Do you really want to give such a primal ascendancy to the word “*the*”? It makes it the head of a head phrase. Hard to imagine that determiners are head words.

Knowing even a little of dependency grammar would have exposed the error: “*the*” should have been D and “*dog*” should have been N (if it stands alone) or NP/D (if it is a noun taking a determiner). But conventional constituency grammars rarely if ever bother with issuing a distinct type for determiners, and thus we arrive at a basic markup error. The road to hell is indeed paved with Chomskian gold.

Type constructors vs. Sexuality

The definition of CCG involved seemingly mono-sexed types, and two type constructors / and \. The definition of LG involves heterosexual types, with connector directions + and – and a single type constructor &. This text has exposed the relationships between these two, but it leaves open a bigger question: what is the formal interplay between type constructors and sexuality? It seems that the one can be traded for the other, but the mechanics of this in a general setting are not clear.

Conclusion

The lesson for today: CCG is equivalent to LG. More or less. We glossed over or completely ignored many of the finer points of LG. No doubt, many important aspects of CCG were omitted as well. Yet, the basic jigsaw structure of CCG was exposed in the plainest way.

The meta-lesson for today: Jigsaws are fundamental for describing a vast class of mathematical and linguistic phenomena. Jigsaws have types (the types of the connectors) and the connectors have “sexuality” (usually heterosexual, for most applications).

The story does not end there; let’s leave off with some hazy futuristic sci-fi: to infinity and beyond! Consider chemical bonds. Two atoms can bond to one-another, using ionic bonds, molecular bonds, hydrogen bonds and van der Waals bonds. In this sense, molecules are clearly jigsaw pieces, having connectors on them. The type+direction theory outlined so far is not quite sufficient to properly describe chemistry. But it does move in that direction. What more is needed to obtain a fully-accurate type-theoretic model of chemistry?

References

- [1] Daniel Sleator and Davy Temperley., *Parsing English with a Link Grammar*, Tech. rep., Carnegie Mellon University Computer Science technical report CMU-CS-91-196, 1991, URL <http://arxiv.org/pdf/cmp-lg/9508004>.